

2_models

December 27, 2025

1 Models training

```
[5]: # --- Set up environment ---
import pandas as pd
import spacy
import numpy as np
import re

from collections import Counter
from collections import defaultdict
from tqdm import tqdm
from nltk.sentiment.vader import SentimentIntensityAnalyzer
from nltk.tokenize import sent_tokenize
from transformers import AutoTokenizer, AutoModelForSequenceClassification
from torch.nn.functional import softmax
from torch import argmax
from pathlib import Path
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
from scipy import sparse

tqdm.pandas()
```

```
[2]: !python -m spacy download en_core_web_sm
nlp = spacy.load("en_core_web_sm")
import features
```

```
Defaulting to user installation because normal site-packages is not writeable
Collecting en-core-web-sm==3.8.0
  Downloading https://github.com/explosion/spacy-
models/releases/download/en_core_web_sm-3.8.0/en_core_web_sm-3.8.0-py3-none-
any.whl (12.8 MB)
----- 0.0/12.8 MB ? eta -:--:--
----- 1.6/12.8 MB 39.8 MB/s eta 0:00:01
----- 12.8/12.8 MB 54.4 MB/s 0:00:00
Download and installation successful
You can now load the package via spacy.load('en_core_web_sm')
```

```
[3]: training_data_pan = "C://Users//andre//OneDrive//Desktop//PAN Task 2025//train//  
    ↴train.jsonl"  
testing_data_pan = "C://Users//andre//OneDrive//Desktop//PAN Task 2025//train//  
    ↴val.jsonl"  
training_data_migration = "C://Users//andre//OneDrive//Desktop//PAN Task 2025//  
    ↴SOCC_data//migration_texts_both_train.xlsx"  
testing_data_migration = "C://Users//andre//OneDrive//Desktop//PAN Task 2025//  
    ↴SOCC_data//migration_texts_both_test.xlsx"
```

2 Baseline model on PAN dataset

```
[8]: # Keep original and metadata  
train_path = Path(training_data_pan)  
val_path = Path(testing_data_pan)  
  
df_train = pd.read_json(train_path, lines=True)  
df_train = df_train[["id", "text", "label", "genre", "model"]].copy()  
df_train["label"] = df_train["label"].astype(int)  
  
# --- Train full baseline model ---  
  
# Apply spaCy to get all tokens, POS tags, etc.  
df_train["spacy_doc"] = df_train["text"].progress_apply(nlp)  
  
# Baseline features for TRAIN  
baseline_df = pd.concat([  
    df_train["spacy_doc"].progress_apply(features.extract_length_features),  
    df_train["spacy_doc"].progress_apply(features.extract_stopword_rate),  
    df_train["text"].progress_apply(features.extract_punct_rates),  
    df_train["text"].progress_apply(features.extract_char_class_ratios),  
    df_train["spacy_doc"].progress_apply(features.extract_ttr),  
    df_train["spacy_doc"].progress_apply(features.extract_upos_freq),  
    df_train["spacy_doc"].progress_apply(features.compute_5gram_repetition),  
    df_train["spacy_doc"].progress_apply(features.compute_self_similarity)  
], axis=1)  
  
# Scale features  
scaler = StandardScaler()  
X_handcrafted_scaled = scaler.fit_transform(baseline_df)  

```

```

df_val = pd.read_json(val_path, lines=True)
df_val["spacy_doc"] = df_val["text"].progress_apply(nlp)

# baseline features for VAL
baseline_val = pd.concat([
    df_val["spacy_doc"].progress_apply(features.extract_length_features),
    df_val["spacy_doc"].progress_apply(features.extract_stopword_rate),
    df_val["text"].progress_apply(features.extract_punct_rates),
    df_val["text"].progress_apply(features.extract_char_class_ratios),
    df_val["spacy_doc"].progress_apply(features.extract_ttr),
    df_val["spacy_doc"].progress_apply(features.extract_upos_freq),
    df_val["spacy_doc"].progress_apply(features.compute_5gram_repetition),
    df_val["spacy_doc"].progress_apply(features.compute_self_similarity)
], axis=1)

# Scale features with SAME scaler
X_handcrafted_scaled_val = scaler.transform(baseline_val)
X_handcrafted_sparse_val = sparse.csr_matrix(X_handcrafted_scaled_val)

# Predict
y_val_pred = clf.predict(X_handcrafted_sparse_val)

df_val = df_val.copy()
df_val["y_true"] = df_val["label"]
df_val["y_pred"] = y_val_pred

# Evaluate
report_df = pd.DataFrame(
    classification_report(
        df_val["label"], y_val_pred,
        target_names=["Human", "AI"],
        output_dict=True
    ).transpose()
)

# Sanity check
print(baseline_df.isna().sum().sum())
print(baseline_val.isna().sum().sum())
print(baseline_df.shape, baseline_val.shape)
print(confusion_matrix(df_val["label"], y_val_pred))

```

```

100%|   | 23707/23707 [32:48<00:00, 12.05it/s]
100%|   | 23707/23707 [00:32<00:00, 732.24it/s]
100%|   | 23707/23707 [00:08<00:00, 2944.71it/s]
100%|   | 23707/23707 [00:05<00:00, 4585.32it/s]
100%|   | 23707/23707 [00:19<00:00, 1222.54it/s]
100%|   | 23707/23707 [00:13<00:00, 1798.40it/s]
100%|   | 23707/23707 [00:10<00:00, 2325.50it/s]

```

```

100%| 23707/23707 [00:24<00:00, 964.43it/s]
100%| 23707/23707 [00:17<00:00, 1371.26it/s]
100%| 3589/3589 [04:49<00:00, 12.40it/s]
100%| 3589/3589 [00:02<00:00, 1235.43it/s]
100%| 3589/3589 [00:01<00:00, 2629.44it/s]
100%| 3589/3589 [00:00<00:00, 6417.69it/s]
100%| 3589/3589 [00:02<00:00, 1289.63it/s]
100%| 3589/3589 [00:01<00:00, 2008.22it/s]
100%| 3589/3589 [00:01<00:00, 2160.12it/s]
100%| 3589/3589 [00:03<00:00, 1060.48it/s]
100%| 3589/3589 [00:02<00:00, 1481.65it/s]

0
0
(23707, 38) (3589, 38)
[[1169 108]
 [ 64 2248]]

```

```
[9]: # --- Analysis ---
# Full report
full_results_output_path = Path("1_baseline_model_pan_results.xlsx")
report_df.to_excel(full_results_output_path, index=True)

print(f"\nSaved baseline model report to: {full_results_output_path}")
print(report_df)

# Misclassifications
misclassified = df_val[df_val["y_true"] != df_val["y_pred"]]
misclassified_feats = baseline_val.loc[misclassified.index]

misclassified_full = pd.concat(
    [misclassified[["id", "text", "y_true", "y_pred", "model", "genre"]],
     misclassified_feats],
    axis=1
)

missclassifications_output_path = Path("baseline_pan_missclassifications.xlsx")
misclassified_full.to_excel(missclassifications_output_path, index=False)

print(f"Saved {len(misclassified_full)} missclassifications to"
      f"{missclassifications_output_path}")

# Check what the baseline classifier learned
handcrafted_df = baseline_df
handcrafted_feat_names = list(handcrafted_df.columns)
coefs = clf.coef_[0]    # shape = (n_features,)

```

```

coef_df = pd.DataFrame({
    "feature": handcrafted_feat_names,
    "weight": coefs
})

top_ai = coef_df.sort_values("weight", ascending=False).head(20)
top_human = coef_df.sort_values("weight", ascending=True).head(20)

print("\nTop features pushing toward AI:")
print(top_ai)

print("\nTop features pushing toward Human:")
print(top_human)

```

Saved baseline model report to: 1_baseline_model_pan_results.xlsx

	precision	recall	f1-score	support
Human	0.948094	0.915427	0.931474	1277.000000
AI	0.954160	0.972318	0.963153	2312.000000
accuracy	0.952076	0.952076	0.952076	0.952076
macro avg	0.951127	0.943873	0.947314	3589.000000
weighted avg	0.952001	0.952076	0.951882	3589.000000

Saved 172 misclassifications to baseline_pan_misclassifications.xlsx

Top features pushing toward AI:

	feature	weight
26	upos_PRON	2.843443
5	punct_>,	1.470360
21	ttr	1.406904
28	upos_DET	1.124349
2	mean_word_len	1.086647
36	rep_5gram_ratio	0.904316
37	self_sim_jaccard	0.888103
0	mean_sent_len	0.652200
19	digit_ratio	0.556874
11	punct_'	0.435686
34	upos_PART	0.391804
3	std_word_len	0.355190
9	punct_?	0.114997
17	upper_ratio	0.057353
30	upos_SCONJ	0.033016
12	punct_"	0.023750
6	punct_.	-0.013307
16	punct_...	-0.022215
8	punct_:	-0.022722
13	punct_(-0.047620

```

Top features pushing toward Human:
      feature      weight
4    stopword_ratio -3.621120
7        punct_ ; -1.504417
35       upos_PROPN -1.285680
15       punct_- -1.240677
1     std_sent_len -1.150756
20      space_ratio -1.017590
31       upos_NUM -0.902325
24       upos_ADJ -0.695302
23       upos_VERB -0.620994
25       upos_ADV -0.478760
33       upos_INTJ -0.463896
32       upos_AUX -0.370075
22       upos_NOUN -0.338835
10      punct_! -0.248612
14      punct_) -0.110865
29       upos_CCONJ -0.094989
27       upos_ADP -0.085844
18      title_ratio -0.062135
13      punct_( -0.047620
8       punct_ : -0.022722

```

3 Baseline model on migration dataset

```
[10]: # Keep original and metadata
train_path = Path(training_data_migration)
val_path = Path(testing_data_migration)

df_train = pd.read_excel(train_path)
df_train = df_train[["id", "text", "label", "genre", "model"]].copy()
df_train["label"] = df_train["label"].astype(int)

# --- Train full baseline model ---

# Apply spaCy to get all tokens, POS tags, etc.
df_train["spacy_doc"] = df_train["text"].progress_apply(nlp)

# Baseline features for TRAIN
baseline_df = pd.concat([
    df_train["spacy_doc"].progress_apply(features.extract_length_features),
    df_train["spacy_doc"].progress_apply(features.extract_stopword_rate),
    df_train["text"].progress_apply(features.extract_punct_rates),
    df_train["text"].progress_apply(features.extract_char_class_ratios),
    df_train["spacy_doc"].progress_apply(features.extract_ttr),
    df_train["spacy_doc"].progress_apply(features.extract_upos_freq),
    df_train["spacy_doc"].progress_apply(features.compute_5gram_repetition),
])
```

```

df_train["spacy_doc"].progress_apply(features.compute_self_similarity)
], axis=1)

# Scale features
scaler = StandardScaler()
X_handcrafted_scaled = scaler.fit_transform(baseline_df)
X_handcrafted_sparse = sparse.csr_matrix(X_handcrafted_scaled)

# Train classifier
clf = LogisticRegression(max_iter=1000, verbose=1)
clf.fit(X_handcrafted_sparse, df_train["label"])

# --- Validate ---
df_val = pd.read_excel(val_path)

df_val["spacy_doc"] = df_val["text"].progress_apply(nlp)

# baseline features for VAL
baseline_val = pd.concat([
    df_val["spacy_doc"].progress_apply(features.extract_length_features),
    df_val["spacy_doc"].progress_apply(features.extract_stopword_rate),
    df_val["text"].progress_apply(features.extract_punct_rates),
    df_val["text"].progress_apply(features.extract_char_class_ratios),
    df_val["spacy_doc"].progress_apply(features.extract_ttr),
    df_val["spacy_doc"].progress_apply(features.extract_upos_freq),
    df_val["spacy_doc"].progress_apply(features.compute_5gram_repetition),
    df_val["spacy_doc"].progress_apply(features.compute_self_similarity)
], axis=1)

# Scale features with SAME scaler
X_handcrafted_scaled_val = scaler.transform(baseline_val)
X_handcrafted_sparse_val = sparse.csr_matrix(X_handcrafted_scaled_val)

# Predict
y_val_pred = clf.predict(X_handcrafted_sparse_val)

# Evaluate
report_df = pd.DataFrame(
    classification_report(
        df_val["label"], y_val_pred,
        target_names=["Human", "AI"],
        output_dict=True
    ).transpose()

# Sanity check results
print(baseline_df.isna().sum().sum())
print(baseline_val.isna().sum().sum())

```

```

print(baseline_df.shape, baseline_val.shape)

print(confusion_matrix(df_val["label"], y_val_pred))

100%| 522/522 [00:47<00:00, 11.09it/s]
100%| 522/522 [00:00<00:00, 712.02it/s]
100%| 522/522 [00:00<00:00, 1994.16it/s]
100%| 522/522 [00:00<00:00, 4790.03it/s]
100%| 522/522 [00:00<00:00, 955.76it/s]
100%| 522/522 [00:00<00:00, 1454.03it/s]
100%| 522/522 [00:00<00:00, 2113.83it/s]
100%| 522/522 [00:00<00:00, 694.19it/s]
100%| 522/522 [00:00<00:00, 833.19it/s]
100%| 78/78 [00:06<00:00, 11.53it/s]
100%| 78/78 [00:00<00:00, 1032.32it/s]
100%| 78/78 [00:00<00:00, 2686.93it/s]
100%| 78/78 [00:00<00:00, 5521.52it/s]
100%| 78/78 [00:00<00:00, 1109.97it/s]
100%| 78/78 [00:00<00:00, 1770.61it/s]
100%| 78/78 [00:00<00:00, 2381.62it/s]
100%| 78/78 [00:00<00:00, 979.61it/s]
100%| 78/78 [00:00<00:00, 1299.25it/s]

0
0
(522, 38) (78, 38)
[[38  0]
 [ 0 40]]

```

```

[11]: # --- Analysis ---
# Full report
output_path = Path("1_baseline_model_migration_db_results.xlsx")
report_df.to_excel(output_path, index=True)

print(f"\nSaved full classification report to: {output_path}")
print(report_df)

# Check what the baseline classifier learned
handcrafted_df = baseline_df
handcrafted_feat_names = list(handcrafted_df.columns)
coefs = clf.coef_[0] # shape = (n_features,)

coef_df = pd.DataFrame({
    "feature": handcrafted_feat_names,
    "weight": coefs
})

```

```

top_ai = coef_df.sort_values("weight", ascending=False).head(20)
top_human = coef_df.sort_values("weight", ascending=True).head(20)

print("\nTop features pushing toward AI:")
print(top_ai)
print("\nTop features pushing toward Human:")
print(top_human)

```

Saved full classification report to: 1_baseline_model_migration_db_results.xlsx

	precision	recall	f1-score	support
Human	1.0	1.0	1.0	38.0
AI	1.0	1.0	1.0	40.0
accuracy	1.0	1.0	1.0	1.0
macro avg	1.0	1.0	1.0	78.0
weighted avg	1.0	1.0	1.0	78.0

Top features pushing toward AI:

	feature	weight
5	punct_ ,	0.657379
22	upos_NOUN	0.626643
3	std_word_len	0.621383
26	upos_PRON	0.582140
2	mean_word_len	0.461726
12	punct_ "	0.426033
28	upos_DET	0.340927
29	upos_CCONJ	0.300897
23	upos_VERB	0.230893
21	ttr	0.206724
24	upos_ADJ	0.134342
8	punct_ :	0.083334
6	punct_ .	0.042621
16	punct_ ...	0.000000
7	punct_ ;	0.000000
10	punct_ !	-0.037683
33	upos_INTJ	-0.060109
9	punct_ ?	-0.067773
25	upos_ADV	-0.085508
36	rep_5gram_ratio	-0.195342

Top features pushing toward Human:

	feature	weight
35	upos_PROPN	-0.735148
15	punct_-	-0.734589
11	punct_ '	-0.727390
18	title_ratio	-0.663105
17	upper_ratio	-0.632454
31	upos_NUM	-0.561319

```

1      std_sent_len -0.509759
32     upos_AUX -0.446126
19     digit_ratio -0.444154
37 self_sim_jaccard -0.405192
20     space_ratio -0.394657
30     upos_SCONJ -0.387993
27     upos_ADP -0.386094
34     upos_PART -0.373883
0      mean_sent_len -0.316033
4      stopword_ratio -0.296349
14     punct_) -0.243967
13     punct_( -0.243967
36 rep_5gram_ratio -0.195342
25     upos_ADV -0.085508

```

4 Bias model on PAN dataset

```
[12]: # Keep original and metadata
train_path = Path(training_data_pan)
val_path = Path(testing_data_pan)

# --- Train bias model ---
df_train = pd.read_json(train_path, lines=True)
df_train = df_train[["id", "text", "label", "genre", "model"]].copy()
df_train["label"] = df_train["label"].astype(int)

# Load lexicons
positive_lexicon = features.load_lexicon(r"C:\Users\andre\OneDrive\Desktop\PAN\u
↪Task 2025\positive-words-dictionary.txt")
negative_lexicon = features.load_lexicon(r"C:\Users\andre\OneDrive\Desktop\PAN\u
↪Task 2025\negative-words-dictionary.txt")
profanity_lexicon = features.load_profanity_lexicon(r"C:
↪\Users\andre\OneDrive\Desktop\PAN Task 2025\toxic-words-dictionary.txt")

# Bias features for TRAIN
bias_df = pd.concat([
    df_train["text"].progress_apply(features.compute_assertive_rate).
↪rename("assertive_rate"),
    df_train["text"].progress_apply(features.compute_sentiment_features),  #_
↪already multi-col
    df_train["text"].progress_apply(lambda t: features.
↪compute_profanity_rate(t, profanity_lexicon)).rename("profanity_rate"),
    df_train["text"].progress_apply(features.compute_subjectivity_score).
↪rename("subjectivity_score"),
    df_train["text"].progress_apply(lambda t: features.compute_hedge_rate(t,_
↪features.hedge_lexicon)).rename("hedge_rate"),
])
```

```

df_train["text"].progress_apply(features.compute_identity_term_rates), # ↵
↪multi-col
df_train["text"].progress_apply(lambda t: features.
↪compute_emotional_tone_from_lexicons(
    t, positive_lexicon, negative_lexicon
)), # multi-col (pos_rate, neg_rate, polarity_score)
df_train["text"].progress_apply(features.compute_categorical_rate).
↪rename("categorical_rate"),
], axis=1)

# Scale features
scaler_bias = StandardScaler()
X_train_bias = scaler_bias.fit_transform(bias_df)

# Train classifier
clf_bias = LogisticRegression(max_iter=1000, verbose=1)
clf_bias.fit(X_train_bias, df_train["label"])

# Bias features for VAL
df_val = pd.read_json(val_path, lines=True)

bias_val = pd.concat([
    df_val["text"].progress_apply(features.compute_assertive_rate).
↪rename("assertive_rate"),
    df_val["text"].progress_apply(features.compute_sentiment_features),
    df_val["text"].progress_apply(lambda t: features.compute_profanity_rate(t, ↵
↪profanity_lexicon)).rename("profanity_rate"),
    df_val["text"].progress_apply(features.compute_subjectivity_score).
↪rename("subjectivity_score"),
    df_val["text"].progress_apply(lambda t: features.compute_hedge_rate(t, ↵
↪features.hedge_lexicon)).rename("hedge_rate"),
    df_val["text"].progress_apply(features.compute_identity_term_rates),
    df_val["text"].progress_apply(lambda t: features.
↪compute_emotional_tone_from_lexicons(
        t, positive_lexicon, negative_lexicon
)),
    df_val["text"].progress_apply(features.compute_categorical_rate).
↪rename("categorical_rate"),
], axis=1)

# Scale using SAME scaler
X_val_bias = scaler_bias.transform(bias_val)

# Predict
y_val_pred = clf_bias.predict(X_val_bias)

```

```

df_val = df_val.copy()
df_val["y_true"] = df_val["label"]
df_val["y_pred_bias"] = y_val_pred

# Evaluate
report_df = pd.DataFrame(
    classification_report(
        df_val["label"], y_val_pred,
        target_names=["Human", "AI"],
        output_dict=True
    )
).transpose()

# Sanity check
print(bias_df.isna().sum().sum())
print(bias_val.isna().sum().sum())
print(bias_df.shape, bias_val.shape)
print(confusion_matrix(df_val["label"], y_val_pred))

```

```

100%| 23707/23707 [00:15<00:00, 1506.26it/s]
100%| 23707/23707 [02:51<00:00, 138.20it/s]
100%| 23707/23707 [00:07<00:00, 3296.95it/s]
100%| 23707/23707 [00:14<00:00, 1671.21it/s]
100%| 23707/23707 [00:08<00:00, 2864.95it/s]
100%| 23707/23707 [00:16<00:00, 1412.47it/s]
100%| 23707/23707 [00:14<00:00, 1659.57it/s]
100%| 23707/23707 [00:16<00:00, 1421.65it/s]
100%| 3589/3589 [00:02<00:00, 1565.04it/s]
100%| 3589/3589 [00:24<00:00, 144.47it/s]
100%| 3589/3589 [00:01<00:00, 2899.33it/s]
100%| 3589/3589 [00:02<00:00, 1696.64it/s]
100%| 3589/3589 [00:01<00:00, 2804.05it/s]
100%| 3589/3589 [00:02<00:00, 1513.17it/s]
100%| 3589/3589 [00:01<00:00, 1944.07it/s]
100%| 3589/3589 [00:02<00:00, 1444.99it/s]

0
0
(23707, 17) (3589, 17)
[[ 967  310]
 [ 252 2060]]

```

```

[13]: # --- Analysis ---
# Full report
full_results_output_path = Path("2_bias_model_pan_results.xlsx")
report_df.to_excel(full_results_output_path, index=True)

```

```

print(f"\nSaved bias model report to: {full_results_output_path}")
print(report_df)

# Misclassifications
misclassified_bias = df_val[df_val["y_true"] != df_val["y_pred_bias"]]
misclassified_bias_feats = bias_val.loc[misclassified_bias.index]

misclassified_bias_full = pd.concat(
    [misclassified_bias[["id", "text", "y_true", "y_pred_bias", "model", "genre"]],
     misclassified_bias_feats],
    axis=1
)

missclassifications_output_path = Path("bias_pan_misclassifications.xlsx")
misclassified_bias_full.to_excel(missclassifications_output_path, index=False)

print(f"Saved {len(misclassified_bias_full)} misclassifications to {missclassifications_output_path}")

# --- Check what the bias classifier learned ---
handcrafted_df = bias_df.copy()
handcrafted_feat_names = list(handcrafted_df.columns)
coefs = clf_bias.coef_[0] # shape = (n_features,)

coef_df = pd.DataFrame({
    "feature": handcrafted_feat_names,
    "weight": coefs
})

top_ai = coef_df.sort_values("weight", ascending=False).head(20)
top_human = coef_df.sort_values("weight", ascending=True).head(20)

print("\nTop bias features pushing toward AI:")
print(top_ai)

print("\nTop bias features pushing toward Human:")
print(top_human)

```

	precision	recall	f1-score	support
Human	0.793273	0.757244	0.774840	1277.00000
AI	0.869198	0.891003	0.879966	2312.00000
accuracy	0.843410	0.843410	0.843410	0.84341
macro avg	0.831236	0.824123	0.827403	3589.00000

```
weighted avg    0.842183  0.843410  0.842561  3589.00000
Saved 562 misclassifications to bias_pan_misclassifications.xlsx
```

Top bias features pushing toward AI:

	feature	weight
14	neg_rate	1.482674
13	pos_rate	1.199948
1	sentiment_mean	0.408142
0	assertive_rate	0.391881
10	orientation_identity_rate	0.133668
11	disability_identity_rate	0.102077
8	religion_identity_rate	0.027694
9	nationality_identity_rate	-0.011181
7	race_ethnicity_identity_rate	-0.195390
15	polarity_score	-0.268695
4	subjectivity_score	-0.299980
12	age_identity_rate	-0.397164
2	sentiment_var	-0.422614
6	gender_identity_rate	-0.475158
5	hedge_rate	-0.537363
3	profanity_rate	-0.684859
16	categorical_rate	-1.162688

Top bias features pushing toward Human:

	feature	weight
16	categorical_rate	-1.162688
3	profanity_rate	-0.684859
5	hedge_rate	-0.537363
6	gender_identity_rate	-0.475158
2	sentiment_var	-0.422614
12	age_identity_rate	-0.397164
4	subjectivity_score	-0.299980
15	polarity_score	-0.268695
7	race_ethnicity_identity_rate	-0.195390
9	nationality_identity_rate	-0.011181
8	religion_identity_rate	0.027694
11	disability_identity_rate	0.102077
10	orientation_identity_rate	0.133668
0	assertive_rate	0.391881
1	sentiment_mean	0.408142
13	pos_rate	1.199948
14	neg_rate	1.482674

5 Bias model on migration dataset

```
[14]: # Keep original and metadata
train_path = Path(training_data_migration)
val_path = Path(testing_data_migration)

df_train = pd.read_excel(train_path)
df_train = df_train[["id", "text", "label", "genre", "model"]].copy()
df_train["label"] = df_train["label"].astype(int)

# --- Train bias model ---

# Load lexicons
positive_lexicon = features.load_lexicon(r"C:\Users\andre\OneDrive\Desktop\PAN\u
↪Task 2025\positive-words-dictionary.txt")
negative_lexicon = features.load_lexicon(r"C:\Users\andre\OneDrive\Desktop\PAN\u
↪Task 2025\negative-words-dictionary.txt")
profanity_lexicon = features.load_profanity_lexicon(r"C:
↪\Users\andre\OneDrive\Desktop\PAN Task 2025\toxic-words-dictionary.txt")

# Bias features for TRAIN
bias_df = pd.concat([
    df_train["text"].progress_apply(features.compute_assertive_rate).
    ↪rename("assertive_rate"),
    df_train["text"].progress_apply(features.compute_sentiment_features),  #✉
    ↪already multi-col
    df_train["text"].progress_apply(lambda t: features.
    ↪compute_profanity_rate(t, profanity_lexicon)).rename("profanity_rate"),
    df_train["text"].progress_apply(features.compute_subjectivity_score).
    ↪rename("subjectivity_score"),
    df_train["text"].progress_apply(lambda t: features.compute_hedge_rate(t,✉
    ↪features.hedge_lexicon)).rename("hedge_rate"),
    df_train["text"].progress_apply(features.compute_identity_term_rates),  #✉
    ↪multi-col
    df_train["text"].progress_apply(lambda t: features.
    ↪compute_emotional_tone_from_lexicons(
        t, positive_lexicon, negative_lexicon
    )),  # multi-col (pos_rate, neg_rate, polarity_score)
    df_train["text"].progress_apply(features.compute_categorical_rate).
    ↪rename("categorical_rate"),
], axis=1)

# Scale features
scaler_bias = StandardScaler()
X_train_bias = scaler_bias.fit_transform(bias_df)
```

```

# Train classifier
clf_bias = LogisticRegression(max_iter=1000, verbose=1)
clf_bias.fit(X_train_bias, df_train["label"])

# --- Bias Feature Classifier Training ---
df_val = pd.read_excel(val_path)

# Bias features for VAL
bias_val = pd.concat([
    df_val["text"].progress_apply(features.compute_assertive_rate).
    ↪rename("assertive_rate"),
    df_val["text"].progress_apply(features.compute_sentiment_features),
    df_val["text"].progress_apply(lambda t: features.compute_profanity_rate(t, ↪
    ↪profanity_lexicon)).rename("profanity_rate"),
    df_val["text"].progress_apply(features.compute_subjectivity_score).
    ↪rename("subjectivity_score"),
    df_val["text"].progress_apply(lambda t: features.compute_hedge_rate(t, ↪
    ↪features.hedge_lexicon)).rename("hedge_rate"),
    df_val["text"].progress_apply(features.compute_identity_term_rates),
    df_val["text"].progress_apply(lambda t: features.
    ↪compute_emotional_tone_from_lexicons(
        t, positive_lexicon, negative_lexicon
    )),
    df_val["text"].progress_apply(features.compute_categorical_rate).
    ↪rename("categorical_rate"),
], axis=1)

# Scale using SAME scaler
X_val_bias = scaler_bias.transform(bias_val)

# Predict
y_val_pred = clf_bias.predict(X_val_bias)

# Evaluate
report_df = pd.DataFrame(
    classification_report(
        df_val["label"], y_val_pred,
        target_names=["Human", "AI"],
        output_dict=True
    )
).transpose()

# Sanity check
print(bias_df.isna().sum().sum())
print(bias_val.isna().sum().sum())
print(bias_df.shape, bias_val.shape)

```

```

print(confusion_matrix(df_val["label"], y_val_pred))

100%|   522/522 [00:00<00:00, 1351.43it/s]
100%|   522/522 [00:03<00:00, 140.72it/s]
100%|   522/522 [00:00<00:00, 3359.57it/s]
100%|   522/522 [00:00<00:00, 1733.85it/s]
100%|   522/522 [00:00<00:00, 2855.02it/s]
100%|   522/522 [00:00<00:00, 1521.81it/s]
100%|   522/522 [00:00<00:00, 2007.11it/s]
100%|   522/522 [00:00<00:00, 1455.84it/s]
100%|   78/78 [00:00<00:00, 1628.84it/s]
100%|   78/78 [00:00<00:00, 145.94it/s]
100%|   78/78 [00:00<00:00, 3290.11it/s]
100%|   78/78 [00:00<00:00, 1706.73it/s]
100%|   78/78 [00:00<00:00, 2828.04it/s]
100%|   78/78 [00:00<00:00, 1484.40it/s]
100%|   78/78 [00:00<00:00, 1960.94it/s]
100%|   78/78 [00:00<00:00, 1469.44it/s]

0
0
(522, 17) (78, 17)
[[36  2]
 [ 1 39]]

```

```

[15]: # --- Analysis ---
# Full report
output_path = Path("2_bias_model_migration_results.xlsx")
report_df.to_excel(output_path, index=True)

print(f"\nSaved bias model report to: {output_path}")
print(report_df)

# Check what the bias classifier learned
handcrafted_df = bias_df.copy()
handcrafted_feat_names = list(handcrafted_df.columns)
coefs = clf_bias.coef_[0]    # shape = (n_features,)

coef_df = pd.DataFrame({
    "feature": handcrafted_feat_names,
    "weight": coefs
})

top_ai = coef_df.sort_values("weight", ascending=False).head(20)
top_human = coef_df.sort_values("weight", ascending=True).head(20)

```

```

print("\nTop bias features pushing toward AI:")
print(top_ai)

print("\nTop bias features pushing toward Human:")
print(top_human)

```

Saved bias model report to: 2_bias_model_migration_results.xlsx

	precision	recall	f1-score	support
Human	0.972973	0.947368	0.960000	38.000000
AI	0.951220	0.975000	0.962963	40.000000
accuracy	0.961538	0.961538	0.961538	0.961538
macro avg	0.962096	0.961184	0.961481	78.000000
weighted avg	0.961817	0.961538	0.961519	78.000000

Top bias features pushing toward AI:

	feature	weight
14	neg_rate	2.183094
13	pos_rate	1.806948
1	sentiment_mean	1.392883
12	age_identity_rate	0.101530
16	categorical_rate	-0.012834
15	polarity_score	-0.040461
5	hedge_rate	-0.076977
2	sentiment_var	-0.082457
4	subjectivity_score	-0.117547
0	assertive_rate	-0.138438
11	disability_identity_rate	-0.325031
8	religion_identity_rate	-0.735063
10	orientation_identity_rate	-0.886660
9	nationality_identity_rate	-1.166516
7	race_ethnicity_identity_rate	-1.577769
3	profanity_rate	-2.122194
6	gender_identity_rate	-2.185796

Top bias features pushing toward Human:

	feature	weight
6	gender_identity_rate	-2.185796
3	profanity_rate	-2.122194
7	race_ethnicity_identity_rate	-1.577769
9	nationality_identity_rate	-1.166516
10	orientation_identity_rate	-0.886660
8	religion_identity_rate	-0.735063
11	disability_identity_rate	-0.325031
0	assertive_rate	-0.138438
4	subjectivity_score	-0.117547
2	sentiment_var	-0.082457
5	hedge_rate	-0.076977

```
15          polarity_score -0.040461  
16          categorical_rate -0.012834  
12          age_identity_rate 0.101530  
1           sentiment_mean 1.392883  
13          pos_rate 1.806948  
14          neg_rate 2.183094
```

6 Merged model on PAN dataset

```

df_train["text"].progress_apply(features.compute_sentiment_features), # ↴
↪already multi-col
df_train["text"].progress_apply(lambda t: features.
↪compute_profanity_rate(t, profanity_lexicon)).rename("profanity_rate"),
df_train["text"].progress_apply(features.compute_subjectivity_score).
↪rename("subjectivity_score"),
df_train["text"].progress_apply(lambda t: features.compute_hedge_rate(t,
↪features.hedge_lexicon)).rename("hedge_rate"),
df_train["text"].progress_apply(features.compute_identity_term_rates), # ↴
↪multi-col
df_train["text"].progress_apply(lambda t: features.
↪compute_emotional_tone_from_lexicons(
    t, positive_lexicon, negative_lexicon
)), # multi-col (pos_rate, neg_rate, polarity_score)
df_train["text"].progress_apply(features.compute_categorical_rate).
↪rename("categorical_rate"),
], axis=1)

# Merge all features
merged_train_full = pd.concat([baseline_train_full, bias_train_full], axis=1)
feature_order = merged_train_full.columns

# Scale
scaler_merged = StandardScaler()
X_train_merged = scaler_merged.fit_transform(merged_train_full)

# Train classifier
clf_merged = LogisticRegression(max_iter=1000, verbose=1)
clf_merged.fit(X_train_merged, df_train["label"])

# --- Validate ---
df_val = pd.read_json(val_path, lines=True)
df_val["spacy_doc"] = df_val["text"].progress_apply(nlp)

# Features from reduced baseline model
baseline_val_full = pd.concat([
    df_val["spacy_doc"].progress_apply(features.extract_length_features),
    df_val["spacy_doc"].progress_apply(features.extract_stopword_rate),
    df_val["text"].progress_apply(features.extract_punct_rates),
    df_val["text"].progress_apply(features.extract_char_class_ratios),
    df_val["spacy_doc"].progress_apply(features.extract_ttr),
    df_val["spacy_doc"].progress_apply(features.extract_upos_freq),
    df_val["spacy_doc"].progress_apply(features.compute_5gram_repetition),
    df_val["spacy_doc"].progress_apply(features.compute_self_similarity)
], axis=1)

```

```

# Bias features
bias_val_full = pd.concat([
    df_val["text"].progress_apply(features.compute_assertive_rate).
    ↪rename("assertive_rate"),
    df_val["text"].progress_apply(features.compute_sentiment_features),
    df_val["text"].progress_apply(lambda t: features.compute_profanity_rate(t, ↪
    ↪profanity_lexicon)).rename("profanity_rate"),
    df_val["text"].progress_apply(features.compute_subjectivity_score).
    ↪rename("subjectivity_score"),
    df_val["text"].progress_apply(lambda t: features.compute_hedge_rate(t, ↪
    ↪features.hedge_lexicon)).rename("hedge_rate"),
    df_val["text"].progress_apply(features.compute_identity_term_rates),
    df_val["text"].progress_apply(lambda t: features.
    ↪compute_emotional_tone_from_lexicons(
        t, positive_lexicon, negative_lexicon
    )),
    df_val["text"].progress_apply(features.compute_categorical_rate).
    ↪rename("categorical_rate"),
], axis=1)

merged_val_full = pd.concat([baseline_val_full, bias_val_full], axis=1)
merged_val_full = merged_val_full.reindex(columns=feature_order)

# Scale using same scaler
X_val_merged = scaler_merged.transform(merged_val_full)

# Predict
y_val_pred = clf_merged.predict(X_val_merged)

df_val = df_val.copy()
df_val["y_true"] = df_val["label"]
df_val["y_pred"] = y_val_pred

# Evaluate
report_df = pd.DataFrame(
    classification_report(
        df_val["label"], y_val_pred,
        target_names=["Human", "AI"],
        output_dict=True
    ).transpose()
)

# Sanity check
print(merged_train_full.isna().sum().sum())
print(merged_val_full.isna().sum().sum())
print(merged_train_full.shape, merged_val_full.shape)

```

```

from sklearn.metrics import confusion_matrix
print(confusion_matrix(df_val["label"], y_val_pred))

100%|   | 23707/23707 [32:44<00:00, 12.07it/s]
100%|   | 23707/23707 [00:32<00:00, 725.36it/s]
100%|   | 23707/23707 [00:09<00:00, 2606.39it/s]
100%|   | 23707/23707 [00:05<00:00, 4184.88it/s]
100%|   | 23707/23707 [00:21<00:00, 1093.65it/s]
100%|   | 23707/23707 [00:16<00:00, 1479.82it/s]
100%|   | 23707/23707 [00:10<00:00, 2181.87it/s]
100%|   | 23707/23707 [00:30<00:00, 772.77it/s]
100%|   | 23707/23707 [00:22<00:00, 1074.84it/s]
100%|   | 23707/23707 [00:15<00:00, 1553.27it/s]
100%|   | 23707/23707 [02:50<00:00, 139.30it/s]
100%|   | 23707/23707 [00:13<00:00, 1703.46it/s]
100%|   | 23707/23707 [00:25<00:00, 941.93it/s]
100%|   | 23707/23707 [00:15<00:00, 1530.98it/s]
100%|   | 23707/23707 [00:29<00:00, 794.24it/s]
100%|   | 23707/23707 [00:21<00:00, 1100.11it/s]
100%|   | 23707/23707 [00:28<00:00, 827.92it/s]
100%|   | 3589/3589 [05:48<00:00, 10.31it/s]
100%|   | 3589/3589 [00:04<00:00, 797.93it/s]
100%|   | 3589/3589 [00:01<00:00, 2355.86it/s]
100%|   | 3589/3589 [00:00<00:00, 5465.54it/s]
100%|   | 3589/3589 [00:03<00:00, 948.46it/s]
100%|   | 3589/3589 [00:02<00:00, 1209.14it/s]
100%|   | 3589/3589 [00:01<00:00, 1968.78it/s]
100%|   | 3589/3589 [00:04<00:00, 750.23it/s]
100%|   | 3589/3589 [00:02<00:00, 1203.76it/s]
100%|   | 3589/3589 [00:02<00:00, 1465.40it/s]
100%|   | 3589/3589 [00:26<00:00, 133.19it/s]
100%|   | 3589/3589 [00:01<00:00, 3282.56it/s]
100%|   | 3589/3589 [00:02<00:00, 1764.81it/s]
100%|   | 3589/3589 [00:01<00:00, 2898.15it/s]
100%|   | 3589/3589 [00:02<00:00, 1282.04it/s]
100%|   | 3589/3589 [00:01<00:00, 1977.84it/s]
100%|   | 3589/3589 [00:02<00:00, 1473.10it/s]

0
0
(23707, 55) (3589, 55)
[[1189  88]
 [ 50 2262]]

```

[17]: # --- Analysis ---
Full report

```

full_results_output_path = Path("3_merged_pan_results.xlsx")
report_df.to_excel(full_results_output_path, index=True)

print(f"\nSaved merged model report to: {full_results_output_path}")
print(report_df)

# Misclassifications
misclassified = df_val[df_val["y_true"] != df_val["y_pred"]]
misclassified_feats = merged_val_full.loc[misclassified.index]

misclassified_merged = pd.concat(
    [misclassified[["id", "text", "y_true", "y_pred", "model", "genre"]],
     misclassified_feats],
    axis=1
)

missclassifications_output_path = Path("merged_pan_missclassifications.xlsx")
misclassified_merged.to_excel(missclassifications_output_path, index=False)

print(f"Saved {len(misclassified_merged)} missclassifications to"
      f"{missclassifications_output_path}")

# Check what the merged classifier learned
handcrafted_df = merged_train_full.copy()
handcrafted_feat_names = list(feature_order)
coefs = clf_merged.coef_[0] # shape = (n_features,)

coef_df = pd.DataFrame({
    "feature": handcrafted_feat_names,
    "weight": coefs
})

# Tag feature origin
coef_df["feature_type"] = np.where(
    coef_df["feature"].isin(baseline_train_full.columns),
    "baseline",
    "bias"
)

top_ai = coef_df.sort_values("weight", ascending=False).head(20)
top_human = coef_df.sort_values("weight", ascending=True).head(20)

print("\nTop merged features pushing toward AI:")
print(top_ai)

print("\nTop merged features pushing toward Human:")
print(top_human)

```

```

Saved merged model report to: 3_merged_pan_results.xlsx
      precision    recall   f1-score   support
Human          0.959645  0.931088  0.945151  1277.000000
AI             0.962553  0.978374  0.970399  2312.000000
accuracy       0.961549  0.961549  0.961549   0.961549
macro avg      0.961099  0.954731  0.957775  3589.000000
weighted avg   0.961518  0.961549  0.961416  3589.000000
Saved 138 misclassifications to merged_pan_misclassifications.xlsx

```

Top merged features pushing toward AI:

	feature	weight	feature_type
26	upos_PRON	2.261831	baseline
2	mean_word_len	1.494952	baseline
5	punct_ ,	1.448293	baseline
21	ttr	1.262385	baseline
28	upos_DET	1.110786	baseline
52	neg_rate	0.999864	bias
37	self_sim_jaccard	0.989597	baseline
36	rep_5gram_ratio	0.833820	baseline
51	pos_rate	0.822653	bias
19	digit_ratio	0.570215	baseline
0	mean_sent_len	0.566931	baseline
11	punct_ '	0.525353	baseline
39	sentiment_mean	0.491009	bias
34	upos_PART	0.380480	baseline
18	title_ratio	0.215660	baseline
9	punct_ ?	0.149644	baseline
17	upper_ratio	0.132821	baseline
3	std_word_len	0.094634	baseline
42	subjectivity_score	0.084035	bias
13	punct_(0.043566	baseline

Top merged features pushing toward Human:

	feature	weight	feature_type
4	stopword_ratio	-2.843400	baseline
7	punct_ ;	-1.579051	baseline
35	upos_PROPN	-1.255422	baseline
15	punct_-	-1.050674	baseline
1	std_sent_len	-1.036793	baseline
20	space_ratio	-1.034956	baseline
24	upos_ADJ	-0.862468	baseline
31	upos_NUM	-0.759919	baseline
25	upos_ADV	-0.564936	baseline
33	upos_INTJ	-0.443324	baseline
23	upos_VERB	-0.436430	baseline
41	profanity_rate	-0.365874	bias
10	punct_ !	-0.305412	baseline

```

32          upos_AUX -0.244425    baseline
53          polarity_score -0.172855   bias
14          punct_) -0.171282    baseline
54          categorical_rate -0.153850   bias
22          upos_NOUN -0.149626    baseline
47 nationality_identity_rate -0.132754   bias
29          upos_CCONJ -0.107629    baseline

```

7 Check a random entry from the PAN data set

```
[23]: import random
import numpy as np
import pandas as pd
from scipy.sparse import csr_matrix
from sklearn.metrics import classification_report

# Select random sample
i = random.randint(0, len(df_val) - 1)
example = df_val.iloc[[i]].copy()
text = example["text"].values[0]

print("Ground truth:", "AI" if example["label"].values[0] == 1 else "Human")
print("\nText excerpt:\n", text[:300], "...")

# Baseline feature extraction
baseline_ex = pd.concat([
    example["spacy_doc"].progress_apply(features.extract_length_features),
    example["spacy_doc"].progress_apply(features.extract_stopword_rate),
    example["text"].progress_apply(features.extract_punct_rates),
    example["text"].progress_apply(features.extract_char_class_ratios),
    example["spacy_doc"].progress_apply(features.extract_ttr),
    example["spacy_doc"].progress_apply(features.extract_upos_freq),
    example["spacy_doc"].progress_apply(features.compute_5gram_repetition),
    example["spacy_doc"].progress_apply(features.compute_self_similarity),
], axis=1)

# scale
X_base_scaled = scaler.transform(baseline_ex)
X_base_sparse = csr_matrix(X_base_scaled)

# predict
proba_base = clf.predict_proba(X_base_sparse)[0][1]
pred_base = clf.predict(X_base_sparse)[0]

# Bias feature extraction
bias_ex = pd.concat([

```

```

example["text"].progress_apply(features.compute_assertive_rate).
    ↪rename("assertive_rate"),
example["text"].progress_apply(features.compute_sentiment_features), # ↴
    ↪already multi-col
example["text"].progress_apply(lambda t: features.compute_profanity_rate(t, ↴
    ↪profanity_lexicon)).rename("profanity_rate"),
example["text"].progress_apply(features.compute_subjectivity_score).
    ↪rename("subjectivity_score"),
example["text"].progress_apply(lambda t: features.compute_hedge_rate(t, ↴
    ↪features.hedge_lexicon)).rename("hedge_rate"),
example["text"].progress_apply(features.compute_identity_term_rates), # ↴
    ↪multi-col
example["text"].progress_apply(lambda t: features.
    ↪compute_emotional_tone_from_lexicons(
        t, positive_lexicon, negative_lexicon
    )), # multi-col (pos_rate, neg_rate, polarity_score)
example["text"].progress_apply(features.compute_categorical_rate).
    ↪rename("categorical_rate"),
], axis=1)

X_bias_scaled = scaler_bias.transform(bias_ex)
proba_bias = clf_bias.predict_proba(X_bias_scaled)[0][1]
pred_bias = clf_bias.predict(X_bias_scaled)[0]

# Merged feature extraction
merged_ex = pd.concat([baseline_ex, bias_ex], axis=1)
merged_ex = merged_ex.reindex(columns=feature_order)

X_merge_scaled = scaler_merged.transform(merged_ex)
proba_merge = clf_merged.predict_proba(X_merge_scaled)[0][1]
pred_merge = clf_merged.predict(X_merge_scaled)[0]

# Results
def show_pred(name, pred, proba):
    print(f"\n{name}: {'AI' if pred==1 else 'Human'}"
          f" (confidence={proba:.3f})")

show_pred("Baseline", pred_base, proba_base)
show_pred("Bias", pred_bias, proba_bias)
show_pred("Merged", pred_merge, proba_merge)

```

Ground truth: AI

Text excerpt:

Brady and Stafford: A Tale of Two Quarterbacks

Los Angeles, CA - In a highly anticipated matchup, two of the NFL's most accomplished quarterbacks, Tom Brady and Matthew Stafford, will face off as the Los Angeles Rams host the Tampa Bay Buccaneers. While their paths to their current teams have been ...

```
100%|   | 1/1 [00:00<00:00, 359.01it/s]
100%|   | 1/1 [00:00<00:00, 1156.09it/s]
100%|   | 1/1 [00:00<00:00, 1264.49it/s]
100%|   | 1/1 [00:00<00:00, 187.98it/s]
100%|   | 1/1 [00:00<00:00, 400.33it/s]
100%|   | 1/1 [00:00<00:00, 406.94it/s]
100%|   | 1/1 [00:00<00:00, 329.04it/s]
100%|   | 1/1 [00:00<00:00, 369.74it/s]
100%|   | 1/1 [00:00<00:00, 390.02it/s]
100%|   | 1/1 [00:00<00:00, 80.59it/s]
100%|   | 1/1 [00:00<00:00, 1546.00it/s]
100%|   | 1/1 [00:00<00:00, 1078.78it/s]
100%|   | 1/1 [00:00<00:00, 1563.29it/s]
100%|   | 1/1 [00:00<00:00, 530.72it/s]
100%|   | 1/1 [00:00<00:00, 658.76it/s]
100%|   | 1/1 [00:00<00:00, 668.20it/s]
```

Baseline: Human (confidence=0.001)

Bias: Human (confidence=0.014)

Merged: AI (confidence=0.583)

8 Test it with your own input

```
[24]: import pandas as pd
import numpy as np
from scipy.sparse import csr_matrix
from sklearn.metrics import classification_report

def predict_and_explain(model_name, clf, feature_df, scaler, ▾
    ↪feature_order=None):
    # If merged model → enforce identical training column order
    if feature_order is not None:
        feature_df = feature_df.reindex(columns=feature_order).fillna(0)
    else:
        feature_df = feature_df.fillna(0)

    # Scale
```

```

X_scaled = scaler.transform(feature_df)

# Sparse for baseline / merged
X_sparse = csr_matrix(X_scaled)

# Predict
proba = clf.predict_proba(X_sparse)[0][1]
pred = clf.predict(X_sparse)[0]
pred_label = "AI" if pred == 1 else "Human"

print(f"\n{model_name} ===")
print(f"Prediction: {pred_label} (confidence={proba:.3f})")

# Explain influence
coefs = clf.coef_[0]
vec = X_sparse.toarray().flatten()
contribs = vec * coefs
feature_names = feature_df.columns.tolist()

top_pos = np.argsort(contribs)[-3:][::-1]
top_neg = np.argsort(contribs)[:3]

print("\nTop features pushing toward AI:")
for idx in top_pos:
    if vec[idx] != 0:
        print(f"{feature_names[idx]:<32} {contribs[idx]:.4f}")

print("\nTop features pushing toward Human:")
for idx in top_neg:
    if vec[idx] != 0:
        print(f"{feature_names[idx]:<32} {contribs[idx]:.4f}")
print("-" * 80)

# Manual input
my_text = """
Time moves differently in places wrapped by mountains and quiet valleys; one
↳day feels as patient as a slow climb, another as fleeting as a bird darting
↳through pines.

In Switzerland, people often walk to the train or stand at a tram stop and
↳glance up at snow-tipped peaks - almost as if the mountains are part of
↳daily rhythm rather than distant scenery.

The multilingual buzz in cafés, from German to French to Italian, carries that
↳subtle sense of multiple worlds coexisting side by side, each greeting the
↳next in its own voice.
"""

```

Everyday order-clean streets, on-time trains, polite commuters-is underpinned ↵by a kind of silent agreement, a willingness to stay small, to notice the ↵detail.

And yet under that calm surface there's movement: ideas flow, debates happen in ↵parliaments and living rooms alike, change sometimes sneaks in as quietly as ↵fog rolling down a slope.

I often wonder if that balance-between stillness and motion, tradition and ↵change-keeps Swiss life alive in a way harder to see on glossy postcards.

Maybe the real charm lies in how ordinary moments carry a little weight, gentle ↵but firm, as steady as the mountains.

"""

```
# Wrap like a one-row dataset
example = pd.DataFrame({"text": [my_text]})
example["spacy_doc"] = example["text"].apply(nlp)

# === Extract each model's feature set exactly like training ===

# Baseline features
baseline_ex = pd.concat([
    example["spacy_doc"].apply(features.extract_length_features),
    example["spacy_doc"].apply(features.extract_stopword_rate),
    example["text"].apply(features.extract_punct_rates),
    example["text"].apply(features.extract_char_class_ratios),
    example["spacy_doc"].apply(features.extract_ttr),
    example["spacy_doc"].apply(features.extract_upos_freq),
    example["spacy_doc"].apply(features.compute_5gram_repetition),
    example["spacy_doc"].apply(features.compute_self_similarity),
], axis=1)

# Bias features
bias_ex = pd.concat([
    example["text"].progress_apply(features.compute_assertive_rate).
    ↵rename("assertive_rate"),
    example["text"].progress_apply(features.compute_sentiment_features), # ↵already multi-col
    example["text"].progress_apply(lambda t: features.compute_profanity_rate(t,
    ↵profanity_lexicon)).rename("profanity_rate"),
    example["text"].progress_apply(features.compute_subjectivity_score).
    ↵rename("subjectivity_score"),
    example["text"].progress_apply(lambda t: features.compute_hedge_rate(t,
    ↵features.hedge_lexicon)).rename("hedge_rate"),
    example["text"].progress_apply(features.compute_identity_term_rates), # ↵multi-col
]
```

```

example["text"].progress_apply(lambda t: features.
    ↪compute_emotional_tone_from_lexicons(
        t, positive_lexicon, negative_lexicon
    )), # multi-col (pos_rate, neg_rate, polarity_score)
example["text"].progress_apply(features.compute_categorical_rate).
    ↪rename("categorical_rate"),
], axis=1)

# Full merged set (same build as training)
merged_ex = pd.concat([baseline_ex, bias_ex], axis=1)

# ===== Run all three models =====
predict_and_explain("BASELINE MODEL", clf, baseline_ex, scaler)
predict_and_explain("BIAS MODEL", clf_bias, bias_ex, scaler_bias)
predict_and_explain("MERGED MODEL", clf_merged, merged_ex, scaler_merged,
    ↪feature_order)

```

100% | 1/1 [00:00<00:00, 2438.55it/s]
100% | 1/1 [00:00<00:00, 350.49it/s]
100% | 1/1 [00:00<00:00, 2824.45it/s]
100% | 1/1 [00:00<00:00, 1952.66it/s]
100% | 1/1 [00:00<00:00, 2641.25it/s]
100% | 1/1 [00:00<00:00, 1186.51it/s]
100% | 1/1 [00:00<00:00, 1399.97it/s]
100% | 1/1 [00:00<00:00, 1974.72it/s]

==== BASELINE MODEL ===

Prediction: AI (confidence=0.721)

Top features pushing toward AI:

upos_AUX	1.6901
std_sent_len	0.9171
self_sim_jaccard	0.7937

Top features pushing toward Human:

upos_PRON	-1.2292
upos_ADP	-1.1505
std_word_len	-0.9150

==== BIAS MODEL ===

Prediction: Human (confidence=0.111)

Top features pushing toward AI:

pos_rate	1.7573
profanity_rate	1.6289

```
sentiment_mean           1.2632

Top features pushing toward Human:
neg_rate                 -4.0221
nationality_identity_rate -3.3572
hedge_rate                -0.1590
-----
-----

==== MERGED MODEL ====
Prediction: Human  (confidence=0.196)

Top features pushing toward AI:
ttr                      4.3853
stopword_ratio            2.5673
std_sent_len               0.9756

Top features pushing toward Human:
upos_PRON                 -3.5327
self_sim_jaccard            -1.8105
upos_ADJ                   -1.6912
-----
```